

ENDSEM SYSTEM PROGRAMMING OPERATING SYSTEM

UNIT – 6

Q.1] Consider page sequence 2, 3, 2, 1, 5, 2, 4, 5, 3, 2, 5, 2 and discuss working of following page replacement policies Also count page faults. (use no. of Frames = 3)

i) FIFO ii) LRU

ANS: Let's analyze the page replacement policies FIFO (First-In-First-Out) and LRU (Least Recently Used) for the given page sequence with a frame size of 3.

Page Sequence: 2, 3, 2, 1, 5, 2, 4, 5, 3, 2, 5, 2

Number of Frames (Frames): 3

1. FIFO (First-In-First-Out):

- **Initialization:**
 - **Frame 1: []**
 - **Frame 2: []**
 - **Frame 3: []**
- **Page Faults:**
 - **Page 2: [2]**
 - **Page 3: [2, 3]**
 - **Page 2: [2, 3]**
 - **Page 1: [1, 3, 2] (Page 2 is replaced as it is the oldest)**
 - **Page 5: [1, 3, 5] (Page 3 is replaced)**
 - **Page 2: [1, 2, 5] (Page 3 is replaced)**
 - **Page 4: [4, 2, 5] (Page 1 is replaced)**
 - **Page 5: [4, 2, 5] (Page 2 is replaced)**
 - **Page 3: [4, 3, 5] (Page 5 is replaced)**
 - **Page 2: [4, 3, 2] (Page 5 is replaced)**
 - **Page 5: [4, 3, 5] (Page 2 is replaced)**
 - **Page 2: [4, 3, 5] (Page 4 is replaced)**

Total Page Faults (FIFO): 9

2. LRU (Least Recently Used):

- **Initialization:**
 - **Frame 1: []**
 - **Frame 2: []**
 - **Frame 3: []**
- **Page Faults:**
 - **Page 2: [2]**
 - **Page 3: [2, 3]**
 - **Page 2: [3, 2]**
 - **Page 1: [1, 3, 2]**
 - **Page 5: [1, 3, 5] (Page 2 is the least recently used and is replaced)**
 - **Page 2: [1, 2, 5]**
 - **Page 4: [4, 2, 5] (Page 1 is replaced)**
 - **Page 5: [4, 2, 5] (Page 3 is replaced)**
 - **Page 3: [4, 3, 5] (Page 2 is replaced)**
 - **Page 2: [4, 3, 2] (Page 5 is replaced)**
 - **Page 5: [4, 3, 5] (Page 2 is replaced)**
 - **Page 2: [4, 3, 2] (Page 4 is replaced)**

Total Page Faults (LRU): 8

In summary, for the given page sequence and a frame size of 3, the LRU page replacement policy results in fewer page faults compared to the FIFO policy. LRU selects pages for replacement based on their least recent usage, while FIFO replaces pages based on their arrival order.

Q.2] Discuss fixed Partitioning and Dynamic Partitioning in detail.

ANS:

Fixed Partitioning:

Definition: Fixed Partitioning is a memory management scheme where the main memory is divided into fixed-size partitions, and each partition can accommodate exactly one process. The number and size of partitions are established during system initialization and remain fixed during the system's operation.

Characteristics:

- 1. Partition Size:**
 - Each partition has a fixed size, and processes are loaded into partitions based on their size.
- 2. No Overlapping:**
 - Partitions do not overlap in the memory. Once a partition is assigned to a process, no other process can use that partition.
- 3. Internal Fragmentation:**
 - Fixed Partitioning suffers from internal fragmentation, where the allocated memory block may be larger than the actual size of the process. This leads to inefficient memory utilization.
- 4. Limited Flexibility:**
 - Since the number and size of partitions are fixed, the system lacks flexibility in accommodating processes of varying sizes.
- 5. Process Placement:**
 - Processes are placed in the partition that can accommodate them, based on their size.

Advantages:

- **Simplicity:** Fixed Partitioning is simple to implement and manage.

Disadvantages:

- **Internal Fragmentation:** Wastage of memory due to internal fragmentation.
- **Limited Flexibility:** The system cannot efficiently handle processes of varying sizes.
- **Inefficient Utilization:** It may lead to inefficient memory utilization.

Dynamic Partitioning:

Definition: Dynamic Partitioning is a memory management scheme where memory is divided into variable-sized partitions to accommodate processes of different sizes. Unlike fixed partitioning, dynamic partitioning allows flexibility in allocating memory blocks according to the size of the process.

Characteristics:

- 1. Variable Partition Size:**
 - Partitions can vary in size, depending on the size of the processes being loaded.
- 2. Overlapping:**
 - Partitions can overlap in memory, allowing more efficient utilization of available memory.
- 3. No Internal Fragmentation:**
 - Dynamic Partitioning eliminates internal fragmentation since the memory allocated to a process precisely matches its size.

4. Process Placement:

- Processes can be placed in any available partition that meets their size requirements.

5. Reclaiming Memory:

- When a process completes, the memory occupied by the process can be reclaimed and used for other processes.

Advantages:

- **Efficient Memory Utilization:** Dynamic Partitioning results in more efficient memory utilization without the internal fragmentation seen in fixed partitioning.
- **Flexibility:** The system can efficiently accommodate processes of various sizes.

Disadvantages:

- **External Fragmentation:** Dynamic Partitioning can suffer from external fragmentation, where free memory is scattered in small, non-contiguous blocks, making it challenging to allocate contiguous memory to larger processes.

Q.3] What is meant by Fragmentation, Explain Buddy Systems Fragmentation in detail?

ANS:

Fragmentation:

Fragmentation refers to the phenomenon where the available memory in a computer system is broken into small, non-contiguous blocks, making it challenging to allocate contiguous blocks of memory to processes. There are two main types of fragmentation:

1. Internal Fragmentation:

- **Internal fragmentation occurs when allocated memory is larger than what is actually needed by a process. As a result, some memory within the allocated block is wasted.**

2. External Fragmentation:

- **External fragmentation occurs when free memory is scattered in small, non-contiguous blocks throughout the system. Even though there might be enough free memory to satisfy a request, it cannot be used if it is not contiguous.**

Buddy System Fragmentation:

The Buddy System is a memory allocation technique that aims to minimize both internal and external fragmentation. In the Buddy System, memory is divided into fixed-size blocks, and each block is a power of 2 in size. The basic idea is to pair up blocks that are powers of 2 in size to form "buddy" pairs.

Here's how the Buddy System works:

1. Initialization:

- **The entire memory is considered one large block, and its size is a power of 2 (e.g., 2^k).**

2. Allocation:

- **When a process requests memory of size S , the system finds the smallest buddy block that is larger than or equal to S . If the block is larger than needed, it is split into two buddies of half its size.**

3. Deallocation:

- **When a block is deallocated, the system checks its buddy. If the buddy is also free, the two buddies are merged back into a larger block. This process continues recursively until the buddy at the original level is found.**

Example:

Let's consider a simplified example with blocks of size 16, 8, 4, and 2:

1. Initialization:

- **Memory is one large block of size 16 (2^4).**

2. Allocation:

- **If a process requests memory of size 4, it is allocated a block of size 4.**
- **If another process requests memory of size 8, it is allocated a block of size 8. The remaining block of size 8 is split into two buddies of size 4 each.**

3. Deallocation:

- **If the process that received a block of size 4 completes, the block is deallocated. The system checks its buddy (if any) and merges them back into a block of size 8.**

This process continues, ensuring that allocated blocks are powers of 2 and minimizing both internal and external fragmentation.

Advantages of Buddy System:

1. Minimizes Internal Fragmentation:

- **Allocating blocks in powers of 2 helps minimize internal fragmentation because the allocated block is always a power of 2.**

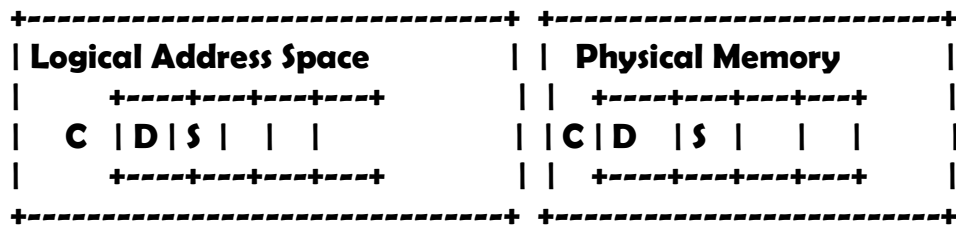
2. Minimizes External Fragmentation:

- **The merging of buddies helps keep the memory contiguous, minimizing external fragmentation.**

3. Efficient Allocation and Deallocation:

- **Allocating and deallocating blocks can be done efficiently by using the buddy system.**

4. Diagram:



In the diagram:

- C represents the Code segment.
- D represents the Data segment.
- \$ represents the Stack segment.

Both Paging and Segmentation are commonly used techniques in modern operating systems to implement virtual memory and enhance the management of memory resources. They offer flexibility, efficient use of memory, and support for larger processes than the physical memory size would allow.

Q.5] Explain Page Table structure and Inverted page Table?

ANS:

Page Table Structure:

A Page Table is a data structure used in virtual memory systems to map virtual addresses to physical addresses. It is an essential component of the memory management unit (MMU) in modern computer architectures. Here's an explanation of the Page Table structure:

1. Page Table Entries (PTE):

- **The Page Table is an array of Page Table Entries, where each entry corresponds to a page or a page frame.**
- **Each entry typically contains information such as the frame number, page attributes (e.g., read-only, read-write), and status bits (e.g., valid or invalid).**

2. Page Table Size:

- **The size of the Page Table is determined by the size of the virtual address space. For example, in a system with a 32-bit virtual address space and 4 KB pages, the Page Table would have 2^{20} entries.**

3. Address Translation:

- **During memory access, the virtual address is divided into a page number and an offset.**
- **The page number is used as an index into the Page Table to obtain the corresponding Page Table Entry.**
- **The frame number from the Page Table Entry is combined with the offset to generate the physical address.**

4. Hierarchical Page Tables:

- **To save memory, large Page Tables can be organized hierarchically. A two-level or multi-level Page Table is used, where the top-level table points to smaller tables.**

5. Page Table Size Consideration:

- **The size of the Page Table affects memory usage. Techniques like multi-level paging or page table compression may be employed to manage large virtual address spaces efficiently.**

6. TLB (Translation Lookaside Buffer):

- **To speed up address translation, a Translation Lookaside Buffer (TLB) is often used. The TLB is a cache of recently used Page Table Entries.**

Inverted Page Table:

An Inverted Page Table is an alternative approach to traditional Page Tables. In this structure, instead of having one entry for each page in the virtual address space, there is one entry for each page frame in physical memory. Here's an explanation of the Inverted Page Table:

1. Structure:

- **The Inverted Page Table is an array with an entry for each frame in physical memory.**
- **Each entry contains the virtual page number, process identifier (PID), and additional information (e.g., status bits).**

2. Address Translation:

- **During memory access, the virtual page number is used to index the Inverted Page Table.**
- **The entry contains information about the mapping, including the process identifier and additional attributes.**

3. Benefits:

- **Reduces memory overhead: Since there is one entry per frame, memory overhead is reduced compared to a conventional Page Table.**
- **Suitable for systems with a large physical memory size.**

4. Hashing:

- **Hashing may be used to locate the corresponding entry quickly. This is important when dealing with a large number of frames.**

5. Efficiency Considerations:

- **Efficient implementation requires careful consideration of collision handling and hashing algorithms to maintain a fast and reliable mapping.**

6. TLB (Translation Lookaside Buffer):

- **Inverted Page Tables may still use a TLB for caching frequently accessed mappings to improve translation speed.**

Comparison:

• Page Table:

- **One entry per virtual page.**
- **Size is proportional to the virtual address space.**
- **Efficient for sparse address spaces.**

• Inverted Page Table:

- **One entry per frame in physical memory.**
- **Size is proportional to the physical memory size.**
- **Efficient for systems with large physical memory sizes.**

Q.6] Given the memory partition of size 100K, 500K, 200K, 300K, 600K, how would each of the First Fit, Best Fit, Worst Fit algorithm place the processes of 212K, 417K, 426K. Which algorithm makes the most efficient use of memory?

ANS: Let's allocate the processes of sizes 212K, 417K, and 426K using the First Fit, Best Fit, and Worst Fit memory allocation algorithms for the given memory partitions (100K, 500K, 200K, 300K, 600K).

Memory Partitions:

1. 100K
2. 500K
3. 200K
4. 300K
5. 600K

Processes:

1. 212K
2. 417K
3. 426K

First Fit Algorithm:

1. **Allocation:**
 - 212K in partition 3 (200K).
 - 417K in partition 5 (600K).
 - 426K does not fit in any available partition.
2. **Efficiency:**
 - **Internal Fragmentation:** $200K - 212K = -12K$ (Excess memory is not used).
 - **External Fragmentation:** $100K + 300K = 400K$ (Unused memory scattered in smaller partitions).

Best Fit Algorithm:

1. **Allocation:**
 - 212K in partition 3 (200K).
 - 417K in partition 5 (600K).
 - 426K in partition 5 (600K) with minimal waste.
2. **Efficiency:**
 - **Internal Fragmentation:** $200K - 212K = -12K$ (Excess memory is not used).
 - **External Fragmentation:** $100K + 300K = 400K$ (Unused memory scattered in smaller partitions).

Worst Fit Algorithm:

1. **Allocation:**
 - 212K in partition 5 (600K).
 - 417K in partition 5 (600K).
 - 426K in partition 5 (600K).
2. **Efficiency:**
 - **Internal Fragmentation:** $600K - 426K = 174K$ (Unused memory in the partition).
 - **External Fragmentation:** $100K + 200K + 300K = 600K$ (Unused memory scattered in smaller partitions).

Efficiency Comparison:

- **First Fit:**
 - High internal fragmentation (negative), high external fragmentation.
- **Best Fit:**

- **High internal fragmentation (negative), high external fragmentation.**
- **Worst Fit:**
 - **High internal fragmentation, moderate external fragmentation.**

In this scenario, none of the algorithms makes particularly efficient use of memory due to significant internal and external fragmentation. However, the Worst Fit algorithm tends to leave larger blocks of unused memory, resulting in less external fragmentation compared to First Fit and Best Fit. The Best Fit algorithm minimizes external fragmentation by selecting the partition with the smallest available space that fits the process. Overall, the efficiency of memory utilization depends on the specific characteristics of the processes and the allocation algorithm used.

Q.7] What is virtual memory management? Explain address translation in paging system.

ANS: here's a simple and concise explanation of virtual memory management and address translation in a paging system:

1. Virtual Memory Management:

- Virtual memory is a technique used by operating systems to provide the illusion of having more physical memory (RAM) than is actually available.
- It allows programs to use more memory than what is physically available by using disk storage as an extension of RAM.

2. Address Translation in Paging System:

- In a paging system, memory is divided into fixed-size blocks called "pages," and processes are divided into blocks of the same size called "page frames."
- When a program references a memory address, it's called a virtual address. This address is divided into two parts: the page number and the offset within the page.
- The operating system maintains a page table that maps virtual page numbers to physical page frame numbers.
- Address translation occurs when the CPU generates a virtual address. The operating system translates this virtual address into a physical address by looking up the page table.
- The page table lookup involves finding the entry corresponding to the virtual page number and retrieving the corresponding physical page frame number.
- Once the physical page frame number is obtained, the offset within the page is added to it to get the final physical address.

3. Key Points:

- Virtual memory management allows efficient use of memory resources by swapping data between RAM and disk.
- Paging simplifies memory management by breaking memory into fixed-size blocks.
- Address translation involves mapping virtual addresses to physical addresses using a page table.
- The page table is maintained by the operating system and contains mappings for virtual pages to physical page frames.
- Address translation enables processes to use memory addresses that may not be physically contiguous.

Q.8] Write proper examples and explain memory allocation strategies first fit, best fit and worst fit. Also explain their advantages and disadvantages.

ANS: here's a breakdown of memory allocation strategies:

1. First Fit:

- **Example:** Imagine you have a memory partition of 100 KB and you need to allocate a process of 60 KB. The first slot you find that is large enough to accommodate the process will be used.
- **Advantages:**
 - Simple and easy to implement.
 - Faster allocation since it searches memory linearly.
- **Disadvantages:**
 - Fragmentation can occur, leading to inefficient memory usage.
 - Suboptimal space utilization since it might pick a larger slot than necessary.

2. Best Fit:

- **Example:** Consider you have memory partitions of 50 KB, 70 KB, and 90 KB. You need to allocate a process of 60 KB. The memory slot closest in size to the process is selected, in this case, the 70 KB partition.
- **Advantages:**
 - Tends to minimize fragmentation by fitting processes into the smallest suitable slot.
 - Better space utilization compared to first fit in some cases.
- **Disadvantages:**
 - More complex implementation compared to first fit.
 - Can be slower since it needs to search for the best fitting slot.

3. Worst Fit:

- **Example:** If you have memory partitions of 30 KB, 80 KB, and 100 KB, and you need to allocate a process of 60 KB, the largest available slot, which is 100 KB, would be selected.
- **Advantages:**
 - Helps to keep larger contiguous free memory spaces.
 - Can lead to less frequent memory allocations, reducing overhead.
- **Disadvantages:**
 - Higher fragmentation compared to best fit.
 - May not be as efficient in space utilization as best fit.

Q.9] Explain any two page replacement strategies in detail.

ANS: Sure, let's break down two popular page replacement strategies:

1. FIFO (First-In-First-Out):

- **Overview:** FIFO replaces the oldest page in memory. It's like waiting in a queue; the first one in is the first one out.
- **Implementation:**
 1. **Maintain a queue to track the order in which pages are brought into memory.**
 2. **When a new page needs to be loaded into memory, check if there's space. If not, remove the page at the front of the queue (the oldest one).**
 3. **Insert the new page at the end of the queue.**
- **Advantages:**
 - **Simple to implement.**
 - **Requires minimal bookkeeping.**
- **Disadvantages:**
 - **Doesn't consider the frequency of page usage or how often a page is accessed.**
 - **Can suffer from the "Belady's Anomaly," where increasing the number of frames can lead to more page faults.**

2. LRU (Least Recently Used):

- **Overview:** LRU replaces the least recently used page. It assumes that the page least recently accessed is the least likely to be needed again soon.
- **Implementation:**
 0. **Maintain a data structure (e.g., a queue, a doubly linked list) to keep track of the order in which pages are accessed.**
 1. **When a new page needs to be loaded into memory:**
 - **If there's space, insert it at the front of the queue (or update its position to the front if it's already present).**
 - **If there's no space, remove the page at the end of the queue (the least recently used one) and insert the new page at the front.**
- **Advantages:**
 - **Generally performs better than FIFO because it considers the recent history of page accesses.**
 - **Less susceptible to Belady's Anomaly.**
- **Disadvantages:**
 - **Requires more complex bookkeeping compared to FIFO.**
 - **Implementing a perfect LRU can be expensive in terms of time and space complexity.**

Q.10] What is TLB? Explain the paging system with the use of TLB? What are the advantages of TLB?

ANS: TLB stands for Translation Lookaside Buffer. Here's a simple explanation of the paging system with the use of TLB and its advantages:

1. Paging System Overview:

- In a paging system, memory is divided into fixed-size blocks called pages.
- Processes are divided into smaller, fixed-size blocks called frames.
- Each process's address space is divided into pages, and each page maps to a frame in memory.

2. Translation Lookaside Buffer (TLB):

- TLB is a small, high-speed cache that stores recently used page table entries.
- It keeps track of the mappings between virtual addresses (used by processes) and physical addresses (used by memory).

3. How TLB Works in Paging:

- When a process accesses memory, the CPU generates a virtual address.
- The virtual address is split into a page number and an offset within the page.
- The TLB is checked first to see if the page number is already in the TLB cache.
- If the page number is found in the TLB, the corresponding physical frame is retrieved directly from memory.
- If the page number is not in the TLB, a page table lookup is performed to find the corresponding physical frame, and the TLB is updated with this mapping.

4. Advantages of TLB:

- **Faster Memory Access:** TLB reduces the time taken to translate virtual addresses to physical addresses by caching frequently accessed mappings.
- **Reduced Memory Access Overhead:** Since TLB stores recently used mappings, it reduces the frequency of accessing the slower page table.
- **Improved Performance:** TLB helps improve overall system performance by speeding up memory access and reducing overhead associated with address translation.
- **Lower Power Consumption:** By reducing the number of memory accesses needed for address translation, TLB can contribute to lower power consumption in the system.

Q.11] What is internal fragmentation? Explain same with suitable diagram/example.

ANS: Sure, I'll break it down for you:

- 1. Definition:** Internal fragmentation occurs when allocated memory is larger than what is actually needed by the process.
- 2. Cause:** It happens due to memory allocation methods where memory is divided into fixed-sized blocks, and if the requested memory is smaller than the block size, the remaining space within the block is wasted.
- 3. Illustration with Diagram:**
 - Let's say we have a memory block divided into fixed-size partitions, like blocks in a puzzle.
 - A process requests memory but doesn't fully utilize the allocated block, leaving some space unused within the block.

[Diagram]

```
|--- Process A (10 units) ---|----- Unused Space (5 units) -----|  
|----- Memory Block (15 units) -----|
```

- 3. Example:**
 - Suppose a program needs 10 units of memory but is allocated a block of 15 units due to fixed-size allocation.
 - After using 10 units, there are 5 units left within the block that can't be used by any other process.
 - This 5-unit unused space within the block is internal fragmentation.
- 4. Consequences:**
 - Reduces overall efficiency of memory utilization.
 - Increases memory wastage over time, especially in systems with frequent allocation and deallocation of memory.
- 5. Mitigation Strategies:**
 - Use dynamic memory allocation techniques to allocate memory based on actual needs.
 - Implement memory compaction algorithms to reclaim fragmented space.
- 6. Importance:**
 - Understanding internal fragmentation is crucial for efficient memory management in operating systems.
 - It influences system performance and resource utilization.
- 7. Conclusion:**
 - Internal fragmentation occurs when allocated memory is larger than required, leading to wastage of memory resources.
 - By implementing appropriate memory management techniques, such as dynamic allocation, internal fragmentation can be minimized, improving overall system efficiency.

Q.12] Write and explain Deadlock Avoidance Bankers Algorithm.

ANS: Here's a simplified explanation of the Banker's Algorithm for deadlock avoidance, broken down into points:

- 1. Introduction:** The Banker's Algorithm is a deadlock avoidance technique used in operating systems to ensure that processes don't get into a deadlock state.
- 2. Resource Allocation:** It involves keeping track of resources (like CPU cycles, memory, etc.) and the maximum resources each process may need to complete its execution.
- 3. Available Resources:** The system maintains a vector called "Available" that represents the number of available resources of each type.
- 4. Maximum Demand:** Each process must declare the maximum number of instances of each resource type that it may need during its execution.
- 5. Allocated Resources:** The system keeps track of the resources currently allocated to each process. This information is stored in a matrix called "Allocation".
- 6. Need Matrix:** The system calculates the resources still needed by each process to complete its execution. This is stored in a matrix called "Need", which is derived from the maximum demand matrix and the allocation matrix.
- 7. Safety Algorithm:** The Banker's Algorithm ensures that there's no deadlock by checking if allocating the requested resources to a process will leave the system in a safe state.
- 8. Safety State:** A state is considered safe if there exists a sequence of processes such that each can obtain all its required resources and then release them, allowing the next process in the sequence to do the same until all processes are done.
- 9. Resource Allocation Process:** When a process requests resources, the Banker's Algorithm checks if granting the request will lead to a safe state. If so, the resources are allocated; if not, the process must wait until its request can be satisfied without causing a deadlock.

Explanation:

- The Banker's Algorithm is like a cautious banker who only grants loans when it's safe to do so. It analyzes the current state of the system and future resource requests to prevent deadlock.
- By keeping track of available resources and the maximum needs of processes, the algorithm ensures that resources are only allocated when it's certain that deadlock won't occur.
- It does this by simulating resource allocation scenarios to see if there's a safe sequence of resource releases, ensuring that processes can complete their execution without getting stuck in a deadlock.
- Overall, the Banker's Algorithm is a proactive approach to managing resources in a way that avoids the possibility of deadlock, promoting system stability and efficiency.

Q.13] Compare Paging and Segmentation with the help of example.

ANS: Here's a comparison between paging and segmentation in a simple and easy-to-understand format:

1. Definition:

- **Paging:** Divides memory into fixed-size blocks called pages. Processes are divided into equal-sized blocks called frames. Allows for efficient memory management but can lead to internal fragmentation.
- **Segmentation:** Divides memory into variable-sized segments based on logical division (such as code, data, stack). Each segment has its own size and can grow dynamically. Can lead to external fragmentation.

2. Memory Allocation:

- **Paging:** Allocates memory in fixed-size blocks (pages) and assigns them to processes.
- **Segmentation:** Allocates memory in variable-sized segments according to the needs of the process.

3. Addressing:

- **Paging:** Addresses are divided into page numbers and page offsets. The page table maps logical addresses to physical addresses.
- **Segmentation:** Addresses are composed of a segment number and an offset within the segment. The segment table maps logical addresses to physical addresses.

4. Fragmentation:

- **Paging:** Internal fragmentation may occur when a page is not fully utilized, leading to wasted memory within a page.
- **Segmentation:** External fragmentation may occur when there are small gaps between segments that are not allocated to any process, leading to inefficient memory usage.

5. Example:

- **Paging:** Imagine a book (memory) divided into equal-sized pages. Each page can contain a different part of the story (process). When you need to access a specific part of the story, you turn to the corresponding page. However, some pages may have extra space at the end, causing internal fragmentation.
- **Segmentation:** Think of a library (memory) where books (processes) are stored in variable-sized sections (segments) based on their genre (code, data, stack). Each book has its own section, and when you want to read a specific book, you go to its corresponding section. However, over time, small gaps may form between sections, leading to external fragmentation.